



(<https://www.udemy.com/user/joseportilla/>).

---

*Copyright by Pierian Data Inc.*

*For more information, visit us at [www.pieriandata.com](http://www.pieriandata.com) (<http://www.pieriandata.com>).*

## KNN Project Exercise - Solutions

Due to the simplicity of KNN for Classification, let's focus on using a PipeLine and a GridSearchCV tool, since these skills can be generalized for any model.

### The Sonar Data

#### Detecting a Rock or a Mine

Sonar (sound navigation ranging) is a technique that uses sound propagation (usually underwater, as in submarine navigation) to navigate, communicate with or detect objects on or under the surface of the water, such as other vessels.



The data set contains the response metrics for 60 separate sonar frequencies sent out against a known mine field (and known rocks). These frequencies are then labeled with the known object they were beaming the sound at (either a rock or a mine).



Our main goal is to create a machine learning model capable of detecting the difference between a rock or a mine based on the response of the 60 separate sonar frequencies.

Data Source: [https://archive.ics.uci.edu/ml/datasets/Connectionist+Bench+\(Sonar,+Mines+vs.+Rocks\)](https://archive.ics.uci.edu/ml/datasets/Connectionist+Bench+(Sonar,+Mines+vs.+Rocks))  
([https://archive.ics.uci.edu/ml/datasets/Connectionist+Bench+\(Sonar,+Mines+vs.+Rocks\)](https://archive.ics.uci.edu/ml/datasets/Connectionist+Bench+(Sonar,+Mines+vs.+Rocks))).

## Complete the Tasks in bold

**TASK: Run the cells below to load the data.**

```
In [95]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [6]: df = pd.read_csv('sonar.all-data.csv')
```

```
In [7]: df.head()
```

Out[7]:

	Freq1	Freq2	Freq3	Freq4	Freq5	Freq6	Freq7	Freq8	Freq9	Freq10	...	Freq52	Freq53	Freq54	Freq55	Freq56	Freq57	Freq58	Fr
0	0.0200	0.0371	0.0428	0.0207	0.0954	0.0986	0.1539	0.1601	0.3109	0.2111	...	0.0027	0.0065	0.0159	0.0072	0.0167	0.0180	0.0084	0.
1	0.0453	0.0523	0.0843	0.0689	0.1183	0.2583	0.2156	0.3481	0.3337	0.2872	...	0.0084	0.0089	0.0048	0.0094	0.0191	0.0140	0.0049	0.
2	0.0262	0.0582	0.1099	0.1083	0.0974	0.2280	0.2431	0.3771	0.5598	0.6194	...	0.0232	0.0166	0.0095	0.0180	0.0244	0.0316	0.0164	0.
3	0.0100	0.0171	0.0623	0.0205	0.0205	0.0368	0.1098	0.1276	0.0598	0.1264	...	0.0121	0.0036	0.0150	0.0085	0.0073	0.0050	0.0044	0.
4	0.0762	0.0666	0.0481	0.0394	0.0590	0.0649	0.1209	0.2467	0.3564	0.4459	...	0.0031	0.0054	0.0105	0.0110	0.0015	0.0072	0.0048	0.

5 rows × 61 columns




## Data Exploration

**TASK: Create a heatmap of the difference frequency responses.**

```
In [ ]: # CODE HERE
```

```
In [8]: plt.figure(figsize=(8,6))  
sns.heatmap(df.corr(),cmap='coolwarm')
```

Out[8]: <AxesSubplot:>



**TASK: What are the top 5 correlated frequencies with the target\label?**

*Note: You may need to map the label to 0s and 1s.*

*Additional Note: We're looking for **absolute** correlation values.*

```
In [99]: #CODE HERE
```

```
In [9]: df['Target'] = df['Label'].map({'R':0, 'M':1})
```

```
In [10]: np.abs(df.corr()['Target']).sort_values().tail(6)
```

```
Out[10]: Freq45    0.339406  
Freq10    0.341142  
Freq49    0.351312  
Freq12    0.392245  
Freq11    0.432855  
Target    1.000000  
Name: Target, dtype: float64
```

## Train | Test Split

Our approach here will be one of using Cross Validation on 90% of the dataset, and then judging our results on a final test set of 10% to evaluate our model.

**TASK: Split the data into features and labels, and then split into a training set and test set, with 90% for Cross-Validation training, and 10% for a final test set.**

```
In [102]: # CODE HERE
```

```
In [11]: from sklearn.model_selection import train_test_split
```

```
In [12]: X = df.drop(['Target', 'Label'], axis=1)  
y = df['Label']
```

```
In [13]: X_cv, X_test, y_cv, y_test = train_test_split(X, y, test_size=0.1, random_state=42)
```

**TASK: Create a Pipeline that contains both a StandardScaler and a KNN model**

```
In [14]: from sklearn.preprocessing import StandardScaler  
from sklearn.neighbors import KNeighborsClassifier
```

```
In [15]: scaler = StandardScaler()
```

```
In [16]: knn = KNeighborsClassifier()
```

```
In [17]: operations = [('scaler',scaler),('knn',knn)]
```

```
In [18]: from sklearn.pipeline import Pipeline
```

```
In [19]: pipe = Pipeline(operations)
```

**TASK: Perform a grid-search with the pipeline to test various values of k and report back the best performing parameters.**

```
In [20]: from sklearn.model_selection import GridSearchCV
```

```
In [21]: k_values = list(range(1,30))
```

```
In [22]: param_grid = {'knn__n_neighbors': k_values}
```

```
In [23]: full_cv_classifier = GridSearchCV(pipe,param_grid,cv=5,scoring='accuracy')
```

```
In [25]: import warnings
warnings.filterwarnings('ignore')

full_cv_classifier.fit(X_cv,y_cv)
```

```
Out[25]: GridSearchCV(cv=5,
                      estimator=Pipeline(steps=[('scaler', StandardScaler()),
                                                  ('knn', KNeighborsClassifier())]),
                      param_grid={'knn__n_neighbors': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11,
                                                       12, 13, 14, 15, 16, 17, 18, 19,
                                                       20, 21, 22, 23, 24, 25, 26, 27,
                                                       28, 29]},
                      scoring='accuracy')
```

```
In [26]: full_cv_classifier.best_estimator_.get_params()
```

```
Out[26]: {'memory': None,
          'steps': [('scaler', StandardScaler()),
                    ('knn', KNeighborsClassifier(n_neighbors=1))],
          'verbose': False,
          'scaler': StandardScaler(),
          'knn': KNeighborsClassifier(n_neighbors=1),
          'scaler__copy': True,
          'scaler__with_mean': True,
          'scaler__with_std': True,
          'knn__algorithm': 'auto',
          'knn__leaf_size': 30,
          'knn__metric': 'minkowski',
          'knn__metric_params': None,
          'knn__n_jobs': None,
          'knn__n_neighbors': 1,
          'knn__p': 2,
          'knn__weights': 'uniform'}
```

**(HARD) TASK:** Using the `.cv_results_` dictionary, see if you can create a plot of the mean test scores per K value.


```
In [113]: #CODE HERE
```

```
In [27]: full_cv_classifier.cv_results_['mean_test_score']
```

```
Out[27]: array([0.84537696, 0.78065434, 0.77524893, 0.75917496, 0.75931721,  
               0.74822191, 0.75945946, 0.71664296, 0.7113798 , 0.68421053,  
               0.70042674, 0.68435277, 0.68449502, 0.67908962, 0.69530583,  
               0.68990043, 0.7113798 , 0.70042674, 0.72204836, 0.67908962,  
               0.70071124, 0.69530583, 0.69530583, 0.68463727, 0.68477952,  
               0.67923186, 0.67411095, 0.65775249, 0.6685633 ])
```

```
In [28]: scores = full_cv_classifier.cv_results_['mean_test_score']  
plt.plot(k_values,scores,'o-')  
plt.xlabel("K")  
plt.ylabel("Accuracy")
```

Out[28]: Text(0, 0.5, 'Accuracy')



## Final Model Evaluation

**TASK:** Using the grid classifier object from the previous step, get a final performance classification report and confusion matrix.



```
In [117]: #Code Here
```

```
In [29]: pred = full_cv_classifier.predict(X_test)
```

```
In [30]: from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
```

```
In [31]: confusion_matrix(y_test, pred)
```

```
Out[31]: array([[12,  1],  
               [ 1,  7]], dtype=int64)
```

```
In [32]: print(classification_report(y_test, pred))
```

	precision	recall	f1-score	support
M	0.92	0.92	0.92	13
R	0.88	0.88	0.88	8
accuracy			0.90	21
macro avg	0.90	0.90	0.90	21
weighted avg	0.90	0.90	0.90	21

**Great Job!**